

1 To get to know the GPU:

```
from numba import cuda
cuda.detect()
```


-> Results:

Found 1 CUDA devices

```
id 0          Tesla T4          [SUPPORTED]
          Compute Capability: 7.5
          PCI Device ID: 4
          PCI Bus ID: 0
          UUID: GPU-0c40aef4-a85b-6723-4274-1db8dde9d5a1
          Watchdog: Disabled
          FP32/FP64 Performance Ratio: 32
```

Summary:

1/1 devices are supported

A screenshot of a Jupyter Notebook interface. The code cell contains the following Python code:

```
from numba import cuda
cuda.detect()
```

 The output cell shows the results of the `cuda.detect()` function. It displays the number of CUDA devices found (1), the device ID (0), the device name (Tesla T4), the compute capability (7.5), the PCI device ID (4), the PCI bus ID (0), the UUID (GPU-0c40aef4-a85b-6723-4274-1db8dde9d5a1), the watchdog status (Disabled), and the FP32/FP64 performance ratio (32). It also includes a summary stating that 1/1 devices are supported. The output ends with `True`. Below the code cell, there is another code cell with

```
cuda.select_device(0)
```

 and its output, which is a weakproxy object to a `numba.cuda.cudadrv.driver.Device` object.

Figure 1: GPU detection result

2 To choose the GPU device

```
cuda.select_device(0)
```

-> Results:

```
<weakproxy at 0x7ad02976f290; to 'numba.cuda.cudadrv.driver.Device' at 0x7ad017b963c0>
```

3 Proof